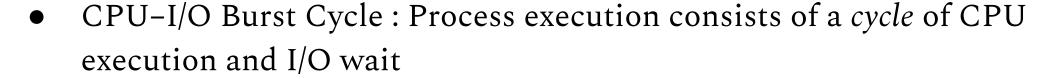
OPERATING SYSTEMS CPU Scheduling

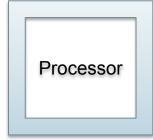
Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- Continuous Cycle:
 - one process has to wait (I/O)
 - Operating system takes the CPU away
 - Give CPU to another process
 - This pattern continues









CPU Scheduler

- Selects from among the processes in ready queue, and allocates the CPU to one of them
 - FIFO queue
 - Priority queue
 - Tree
 - Unordered linked-list
- CPU scheduling decisions may take place when a process:
 - 1. Switches from running to waiting state (I/O request)
 - 2. Switches from running to ready state (e.g. when interrupt occurs)
 - 3. Switches from waiting to ready (e.g. at completion of I/O)
 - 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is *preemptive*
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities

Scheduling Criteria

- **CPU utilization** keep the CPU as busy as possible
- Throughput # of processes that complete their execution per time unit

• Turnaround time

- amount of time to execute a particular process
- -- the interval from the time of submission of a process to the time of the completion.
- -- sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, doing I/O
- Waiting time amount of time a process has been waiting in the ready queue
- **Response time** amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

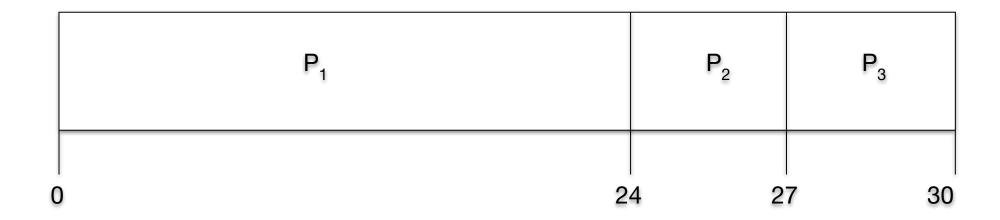
OPERATING SYSTEMS

CPU Scheduling Algorithms - First Come First Serve (FCFS)

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
P_{1}	24
P_{2}	3
P_{3}	3

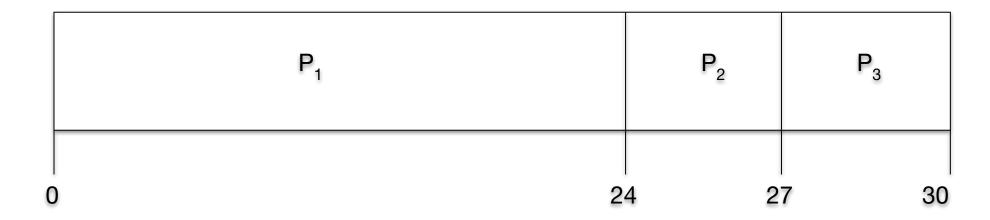
• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is:



First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
P_{1}	24
P_{2}	3
P_{3}	3

• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is:



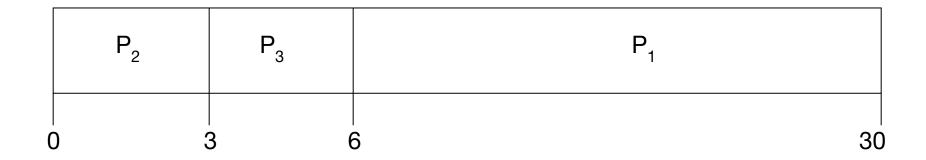
- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: (0 + 24 + 27)/3 = 17
- Turnaround time $P_1 = 24$; $P_2 = 27$; $P_3 = 30$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

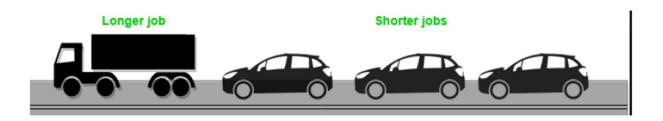
• The Gantt chart for the schedule is:



Waiting time for $P_1 = 6$; $P_2 = 0$, $P_3 = 3$ Average waiting time: (6 + 0 + 3)/3 = 3Much better than previous case

Convoy effect - short process behind long process.

Consider one CPU-bound and many I/O-bound processes



OPERATING SYSTEMS

CPU Scheduling Algorithms - Shortest Job First (SJF)

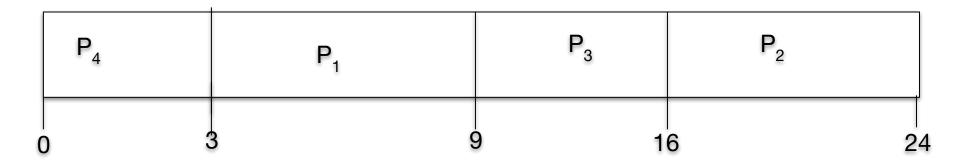
Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - Non-preemptive once CPU given to the process it cannot be preempted until completes its CPU burst
 - preemptive if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal gives minimum average waiting time for a given set of processes

Example of SJF

ProcessBurst Time P_1 6 P_2 8 P_3 7 P_4 3

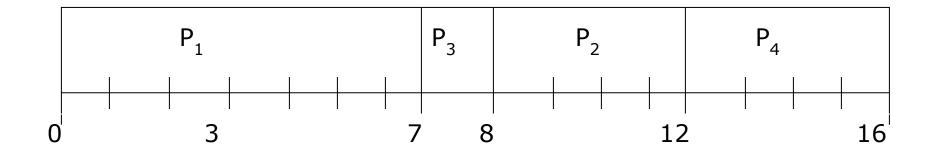
• SJF scheduling chart



Example of Non-Preemptive SJF

<u>Process</u>	<u> Arrival Time</u>		<u>Burst Time</u>
P_{1}	0.0	7	
P_{2}	2.0	4	
P_3	4.0	1	
$P_{_{4}}$	5.0	4	

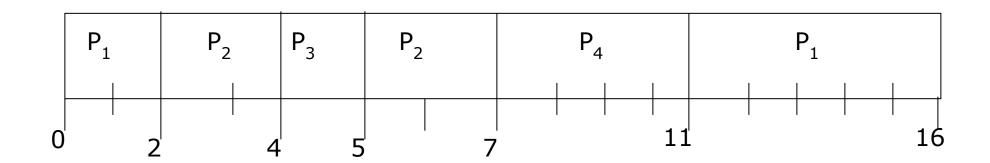
• SJF (non-preemptive)



Example of Preemptive SJF

<u>Process</u>	<u> Arrival Time</u>		Burst Time
P_{1}	0.0	7	
$P_{2}^{}$	2.0	4	
P_{3}	4.0	1	
$P_{_{4}}$	5.0	4	

• SJF (preemptive)

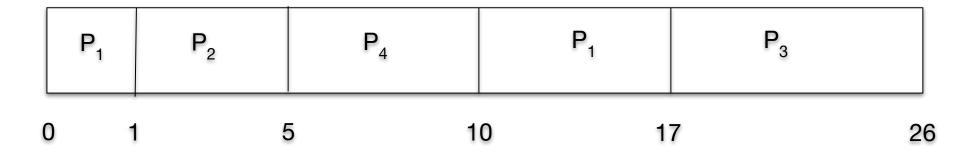


Example of Shortest-remaining-time-first

• Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u> Arrival Time</u>	<u>Burst Time</u>
P1	0	8
P2	1	4
Р3	2	9
P4	3	5

• Preemptive SJF Gantt Chart



OPERATING SYSTEMS

CPU Scheduling Algorithms - Priority Scheduling

Priority Scheduling

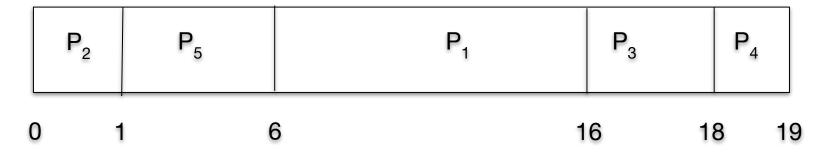
- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer ≡ highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Priority can be defined either internally or externally.
 - Factors for internal priority assignment:
 - Time limit, memory requirements, the number or open files etc.
 - Factors for external priority assignment:
 - Importance of the process, the type and amount of funds being paid for computer use, department sponsoring works etc.

Example of Priority Scheduling

Non Preemptive:

<u>Process</u>	<u>Bu</u> :	<u>rst Time</u>	<u>Priority</u>
$P_{_{1}}$	10	3	
P_{2}	1	1	
$P_{\overline{3}}$	2	4	
$P_{_{4}}$	1	5	
$P_{\overline{5}}$	5	2	

Priority scheduling Gantt Chart



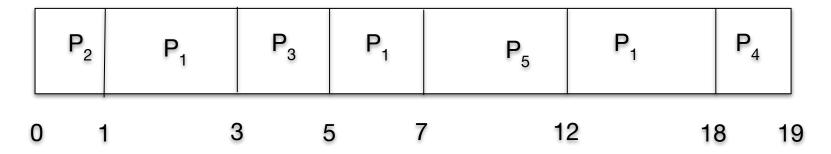
- Problem ≡ **Starvation** low priority processes may never execute
- Solution \equiv **Aging** as time progresses increase the priority of the process

Example of Priority Scheduling

Preemptive:

<u>Process Arrival Time Burst Time Prior</u>	<u> Ity</u>
$P_1 0 \qquad \qquad 10 \qquad \qquad 4$	
$P_2 0 \qquad \qquad 1 \qquad \qquad 1$	
P_{3} 3 2 3	
P_4 5 1 5	
P_{5} 7 5 2	

Priority scheduling Gantt Chart



- Problem ≡ **Starvation** low priority processes may never execute
- Solution \equiv **Aging** as time progresses increase the priority of the process

OPERATING SYSTEMS

CPU Scheduling Algorithms - Round Robin (RR)

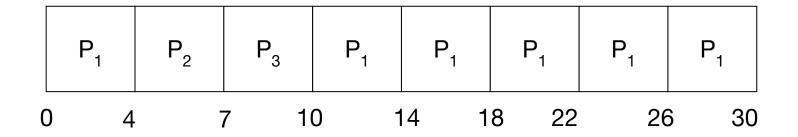
Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q, then each process gets 1/n of the CPU time in chunks of at most q time units at once. No process waits more than (n-1)q time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - \circ q large \Rightarrow FIFO
 - o $q \text{ small} \Rightarrow q \text{ must be large with respect to context switch, otherwise overhead is too high}$

Example of RR with Time Quantum = 4

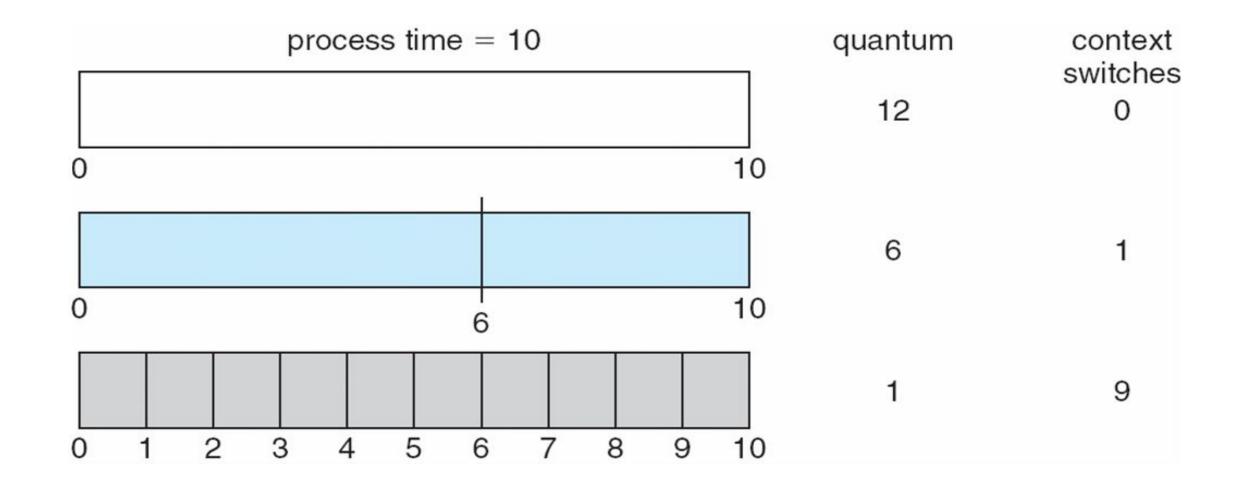
<u>Process</u>	Burst Time
$P_{_1}$	24
\vec{P}_2	3
P_{3}^{2}	3
J	

The Gantt chart is:

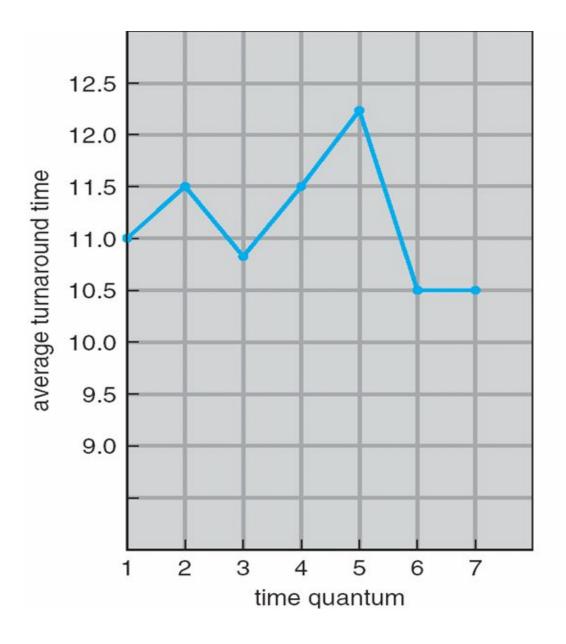


- Average waiting time is 17/3 = 5.66 milisecond
- Typically, higher average turnaround than SJF, but better response
- quantum should be large compared to context switch time
- Quantum usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts should be shorter than quantum

OPERATING SYSTEMS

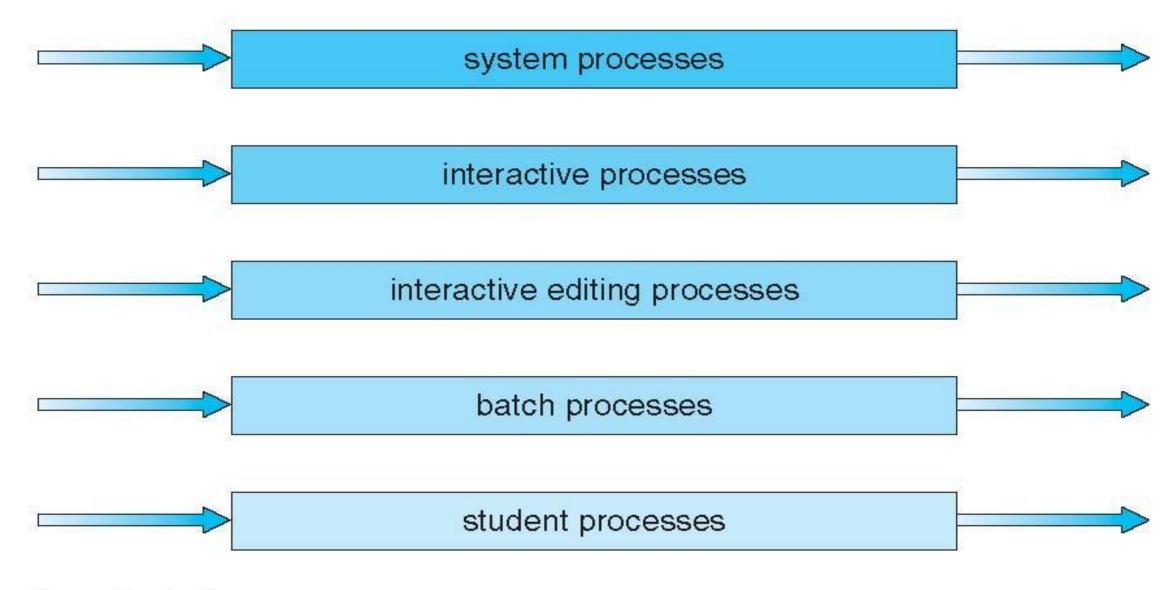
CPU Scheduling Multilevel Queue, Multilevel Feedback Queue

Multilevel Queue

- Another class of scheduling algorithm needs- in which processes are classified into different groups, e.g.:
 - foreground (interactive) processes
 - background (batch) processes
- They have different response time requirements-so different scheduling needs.
- Foreground processes may have priority over background processes.
- A multilevel queue-scheduling algorithm partitions the ready queue into several separate queues-we
 can see it in the figure of next slide:-
- Each queue has its own scheduling algorithm:
 - Foreground queue scheduled by RR algorithm
 - Background queue scheduled by FCFS algorithm
- Scheduling must be done between the queues:
 - Fixed priority preemptive scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., foreground queue can be given 80% of the CPU time for RR-scheduling among its processes, while 20% to background in FCFS manner.

Multilevel Queue Scheduling

highest priority



lowest priority

Multilevel Feedback Queue scheduling

- Multilevel Feedback Queue scheduling, allows a process to move between queues.
- If a process uses too much CPU time, it will be moved to a lower priority queue.
- Similarly, a process that waits too long in a lower-priority queue may me moved to a higher-priority queue.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - o scheduling algorithms for each queue
 - o method used to determine when to upgrade a process
 - o method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

- Three queues: (can see the figure in next slide)
 - \circ Q_0 RR with time quantum 8 milliseconds
 - \circ Q_1 RR time quantum 16 milliseconds
 - \circ $Q_2 FCFS$
- Scheduling
 - A new job enters queue Q_0 which is served for RR
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q₁
 - At Q₁ job is again served RR and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q₂

