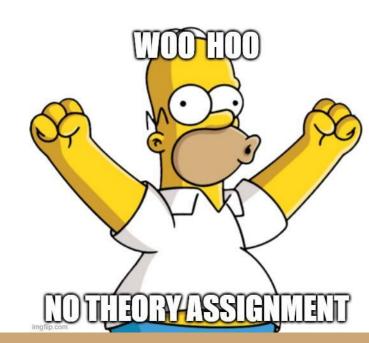
# CSE321: Operating Systems Introduction

#### **Course Outcome**

- To understand the fundamental concepts of computer system organization and the structure of operating systems.
- To explore various aspects of process management in operating system
- To know how different CPU scheduling algorithm works and their respective importance
- To develop practical knowledge on the concept of threads
- To inspect process synchronization mechanisms and deadlocks
- To be able to analyze the management of main and virtual memory

#### **Marks Distribution**

- Theory 80%
  - Class participation 5%
  - 。 Quiz 15%
  - ∘ Mid 25%
  - ∘ Final 35%
- Lab 20%



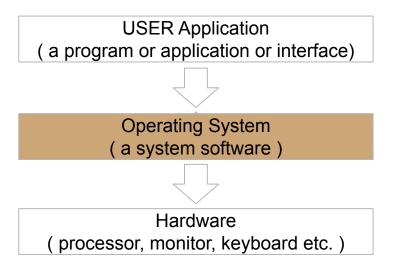
#### **Operating Systems**

#### "Actual" Introduction



# What is an Operating System?

A program that acts as an intermediary between a user of a computer and the computer hardware.



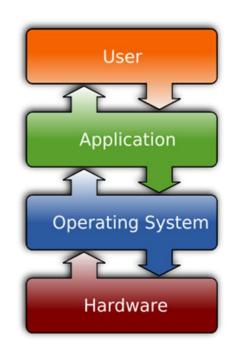
# System Software Vs Application Software

#### **System Software:**

- System Software refers to the operating system and all utility programs that manage computer resources at a low level.
- Systems software includes compilers, loaders, linkers, and debuggers.

#### **Application Software:**

 Applications software comprises programs designed for an end user, such as word processors, database systems, and spreadsheet programs.



#### **Major Goals of OS**

- Execute user programs.
- Make the computer system convenient to use.
- Use the computer hardware in an efficient manner
- Manages and allocate all resources
- Controls the execution of user programs and operations of I/O devices

GM-NAA I/O, produced by General Motors for its **IBM 704** 

1956



MS-DOS is Apple ][ released by released Microsoft

1977 1981



Linux is released by Linus **Torvalds** 

1991

is released 1995

Microsoft Windows 95

Windows 95



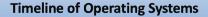
Android is released (based on a Linux kernel)

OpenShift released by Red Hat

2008

2011

2010s





IBM develops a series of OSs for its 360 series. Multics is developed and abandoned but **UNIX** is developed as a consequence.



Unix becomes popular in academic circles and spawns many versions











The home computer revolution



#### 1990s

Windows dominates the laptop and desktop market



Unix and then Linux dominate the Supercomputer Market



#### 2000s

**Smart phones** become ubiquitous after the iPhone release in 2007

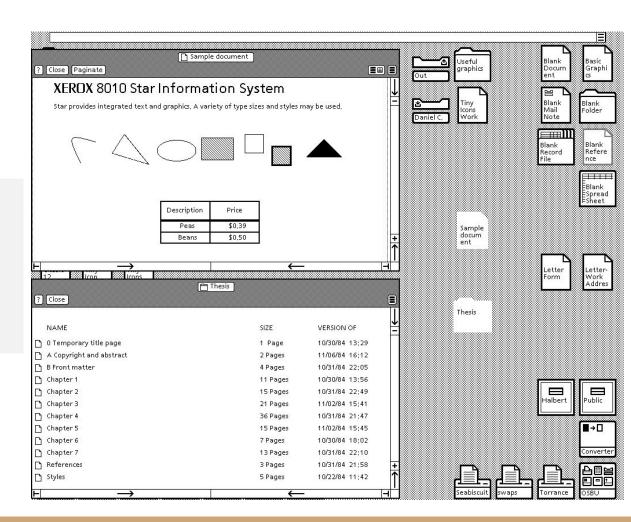




- IBM 704 was the first mass-produced computer
- GM-NAA (General Motors-North American Aviation) I/O is the OS used in the IBM 704
- There was no UI in the OS as we know now
- It was mostly based on command prompts (terminals)



 Xerox 8010 Star (released in 1981) was the first system that was referred to as a fully integrated desktop computer including applications and a GUI



- Bill Gates got lucky to sign the so-called "Deal of the Century" with IBM to design a new OS for IBM machines
- The deal allowed Microsoft to add an OS with \$50 per PC and IBM did not own copyright for the OS
- This allowed Microsoft to start the their OS dominance on the PC domain

- Developed by Apple Computer, Inc for their new product, the Macintosh home PC, The Macintosh 128K, was released in 1984
- was widely advertised (the famous 1984 commercial is available below).
- Mac OS was the first OS with a GUI built-in
- It was also one of the first consumer computers with mouse!
- Even though Microsoft introduced mouse in their PCs a bit earlier
- The use of GUI with mouse was not Steve Job's idea, the idea was taken during his visit in Xerox PARC (Palo Aalto Research Center)
- Jobs reportedly traded US \$1 million in stock options to Xerox for a detailed tour of their facilities and current projects
- One of the things Xerox showed Jobs was the Alto, which sported a GUI and a three-button mouse









- Another revolution of OS came in the mobile computing domain, when Steve Jobs introduced iPhone with iOS in 2007
- The iPhone introduction video is now regarded as a classic advertise video

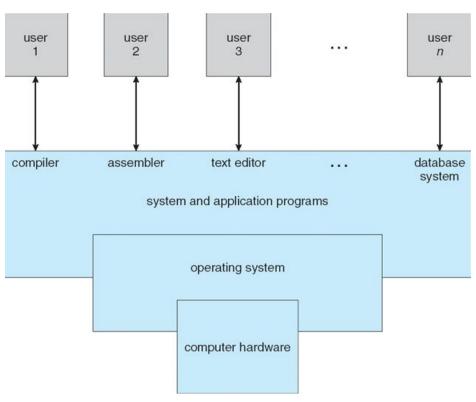


# AN UPDATE IS AVAILABLE FOR YOUR COMPUTER



# Operating Systems Computer System Organization

# **Components of a Computer System**



#### Kernel

The one program running at all times.

- Kernel is the central module of an operating system
- Part of OS that loads first, and it remains in main memory.
- As small as possible
- Provide all the essential services required by other parts of the operating system and applications.
- Kernel code is usually loaded into a protected area of memory to prevent it from being overwritten.

#### **Bootstrap Program**

An initial program executed when a computer starts running.

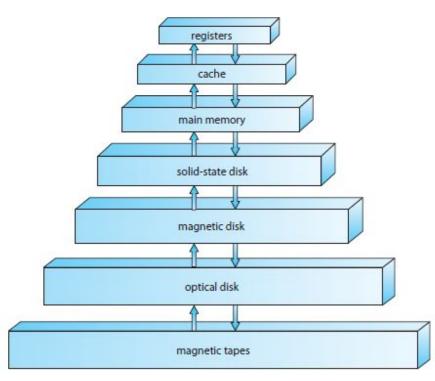
- When a computers is powered up or rebooted, it is executed first.
- Stored in the ROM or EEPROM, known as firmware
- Initializes all aspects of the system, from CPU registers to device controllers to memory contents
- bootstrap program must know how to load the operating system and how to start executing
- Once the OS kernel is loaded and executing, it can start providing services to the system and its users

#### **Storage Structure**

- Main memory only large storage media that the CPU can access directly
  - Random access
  - Typically volatile
- □ Secondary storage extension of main memory that provides large nonvolatile storage capacity

- CPU can load instructions only from main memory.
- General-purpose computers run most of their programs from rewritable memory, called main memory ( also called RAM)
- Computers use other forms of memory as well Read Only Memory (ROM) and electrically erasable programmable read-only memory (EEPROM)
- Only static programs, such as the bootstrap program described earlier, are stored here.

# **Storage Device Hierarchy**

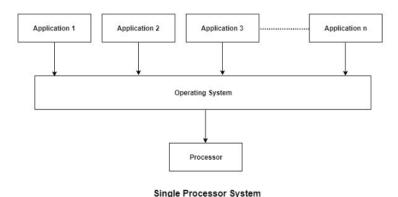


# Operating Systems OS Architecture

#### **Operating System Architecture**

#### **Single-Processor Systems:**

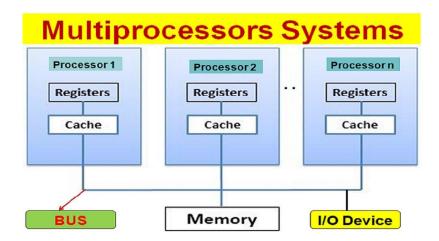
- One main CPU capable of executing a general-purpose instruction set.
- Almost all single processor systems have other special-purpose processors (device-specific processors), which run a limiter instruction set.

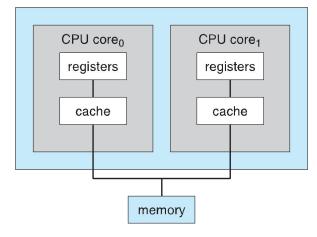


## **Operating System Architecture**

#### Multiprocessor Systems (parallel systems or multicore systems ):

• Such systems have two or more processors in close communication, sharing the computer bus and sometimes the clock, memory, and peripheral devices.



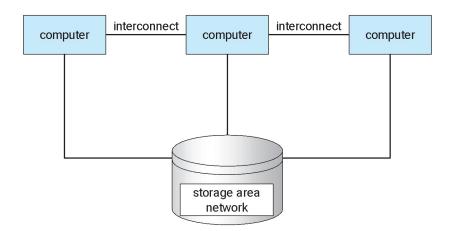


A dual-core system

#### **Operating System Architecture**

#### **Clustered Systems:**

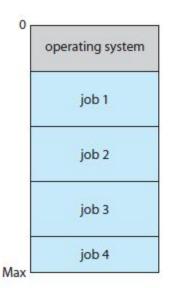
- Special kind of multiprocessor system which gathers together multiple CPUs. They are composed of two or more individual systems- or nodes - joined together.
- Clustered computers share storage and are closely linked via a local-area network (LAN) or a faster interconnect, such as InfiniBand



#### **Operating System Structure**

#### **Multiprogramming:**

- Increases CPU utilization by organizing jobs (code and data) so that the CPU always has one to execute.
- OS keeps several jobs in memory simultaneously. As main memory is small, it keeps the jobs on the disk (job pool) which waits there to be allocated in the main memory.
- OS picks and begins to execute one of the jobs in memory.
   Eventually, the job may have to wait for some task. OS picks another job from the pool to execute while the previous one waits.



Memory layout for a multiprogramming system.

## Requirements of Multiprogramming

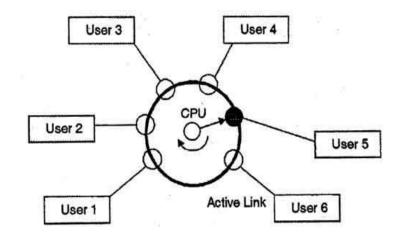
- **Job Scheduling:** If several jobs are ready to be brought into memory, and if there is not enough room for all of them, then the system must choose among them.
- When OS selects a job from the job pool, it loads that job into memory for execution. Having several programs in memory at the same time requires some form of memory management.
- CPU Scheduling: if several jobs are ready to run at the same time, the system must choose which job will run first.
- Running multiple jobs concurrently requires that their ability to affect one another be limited in all phases of the operating system.

★ If processes don't fit in memory, swapping moves them in and out to run from main memory achieving this goal is virtual memory.

## **Operating System Structure**

#### **Time Sharing:**

- CPU executes multiple jobs by switching among them.
- Switches occur so frequently that the users can interact with each program while it is running.
- requires an interactive computer system, which provides direct communication between the user and the system.
- Response time should be short.

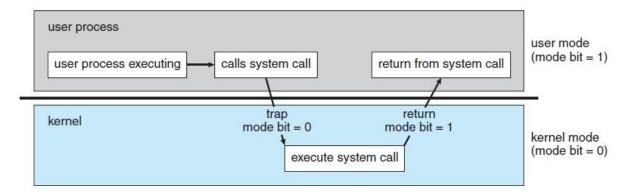


#### **Operating System Operations**

- Modern operating systems are interrupt driven
- Events are almost always signaled by the occurrence of an interrupt or a trap
- A trap (or an exception) is a software-generated interrupt
- For each type of interrupt, separate segments of code determine what action should be taken
- Errors can occur when one erroneous program modify another program, the data of another program, or even the operating system itself.
- A properly designed operating system must ensure that an incorrect (or malicious) program cannot cause other programs to execute incorrectly.

## **Dual Mode Operation**

- Need to distinguish between the execution of operating-system code and user defined code.
- A bit, called the mode bit, is added to the hardware of the computer to indicate the current mode: kernel (0) or user (1).
- dual mode of operation provides protection of the operating system from errant users
- this protection is provided by designating some of the machine instructions that may cause harm as privileged instructions that are executed only in kernel mode.

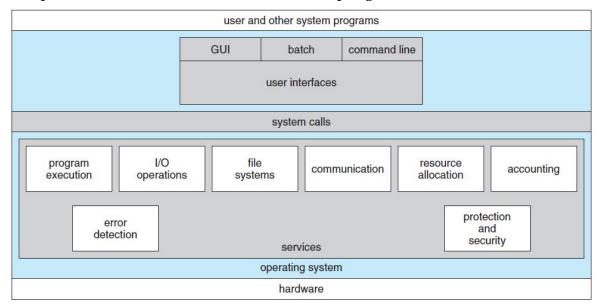


Transition from user to kernel mode.

# Operating Systems OS Services

# **Operating System Services**

- OS provides an environment for the execution of programs.
- Specific services provided, differ from one operating system to another, but there are some common classes
- Services are provided for the convenience of the programmer

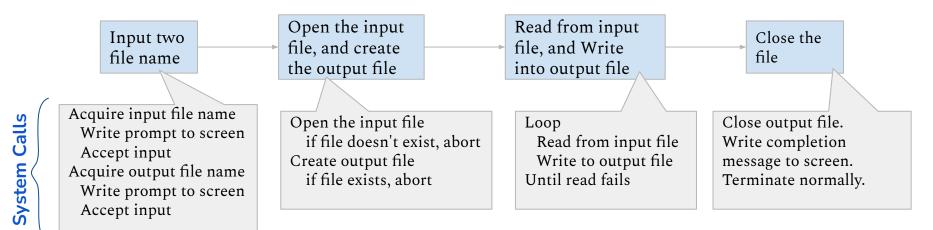


Operating Systems
System Call,
System Program,
System Boot

# System Call

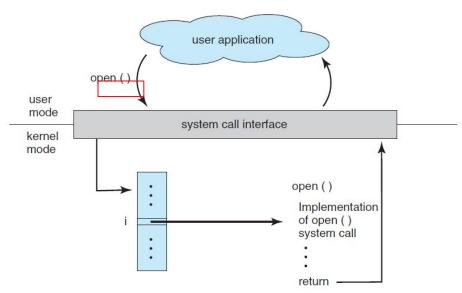
- System calls provide an interface to the services made available by an operating system.
- These calls are generally available as routines.
- Routines are written in C or C++. Some low level tasks are written in assembly language.

A program to copy the contents of a file to another file!



# System Call Interface

- Serves as the link to system calls made available by the operating system.
- A number is associated with each system call, and the system-call interface maintains a table indexed according to these numbers.
- Invokes the intended system call in the operating-system kernel and returns the status of the system call and any return values.



# Types of System Call

Туре	Windows OS	Linux OS
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()

# Types of System Call

Type	Windows OS	Linux OS
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitlializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

## System Programs

System programs, also known as system utilities, provide a convenient environment for program development and execution.

These system programs provide -

- File management
- Status information
- File modification
- Programming-language support
- Program loading and execution
- Communications
- Background services

## System Boot

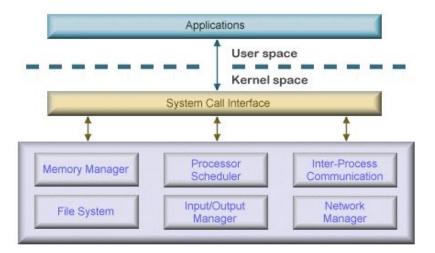
- When power initialized on system, execution starts at a fixed memory location.
  - → Firmware ROM used to hold initial boot code
- Operating system must be made available to hardware so hardware can start it.
  - → Small piece of code bootstrap loader, stored in ROM locates the kernel, loads it into memory, and starts it
  - → Sometimes two-step process where boot block at fixed location loaded by ROM code, which loads bootstrap loader from disk
- Common bootstrap loader, GRUB, allows selection of kernel from multiple disks, versions, kernel options
- Kernel loads and system is then running.

# Operating Systems OS Structures

#### OS Structure

#### Simple/Monolithic structure:

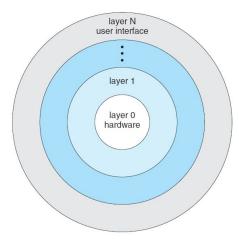
- Earliest and most common architecture
- Every component of OS is in the kernel and can communicate with each other directly
- Complex and Large (millions line of code,) hard to maintain



#### **OS Structure**

#### Layered structure:

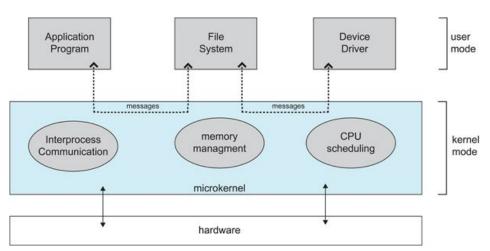
- OS is divided into layers.
- Each layer can use services of its lower layers.
- Easy to debug and develop.
- Less efficient as each layer adds some overhead



#### **OS Structure**

#### Microkernel structure:

- Moves as much from kernel into user space
- Communication takes place between user modules using message passing
- Easier to extend a microkernel
- More reliable( less code running in kernel mode ) and secure
- Performance overhead





## Upcoming Episodes

#### **Process:**

- Process States: A condition of the process at a specific instant of time.
- Process Architecture: The structural design of general process systems.
- Operations on Process: process creation, preemption, blocking, and termination etc
- Process Scheduling: A task that schedules processes of different states like ready, waiting, and running
- Process Synchronization: The task of coordinating the execution of processes

## **Upcoming Episodes**

#### Thread:

- **Definition & Purpose:** a sequential flow of tasks within a process.
- Multithreading Models: Ways of achieving multithreading
- Thread Library: provides the programmer with an Application program interface for creating and managing thread
- Threading Issues: Common problems and pitfalls with multi-thread programming

# Upcoming Episodes

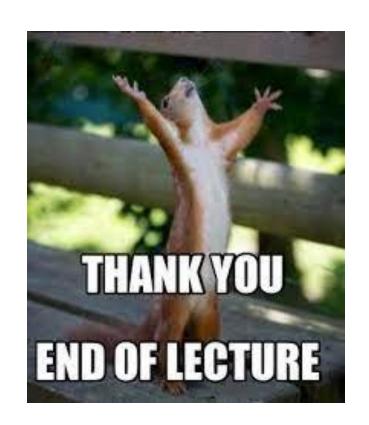
#### **Deadlock:**

- Reason
- Necessary Conditions for Deadlock
- Methods of Handling Deadlocks



# Upcoming Episodes Memory Management:

- Background and importance of memory management
- Memory Management Techniques
- Virtual Memory



# Operating Systems **Process**

## Process Concept

What to call the activities of CPU?

Jobs

Batch System

User Programs or Tasks

Time Sharing
System

These activities are called "Processes"

★ The terms "job" and "process" are used almost interchangeably.

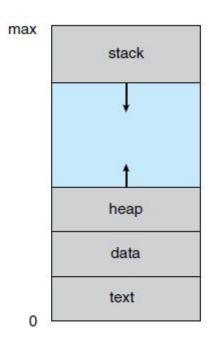
#### Process

A process is a program that is in execution.

But, it is more than the program codes. Program code is known as "text section" of a process.

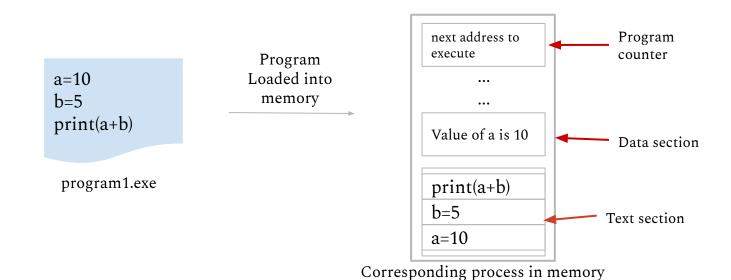
Besides code of the program, it contains -

- Program Counter and Registers: stores current activity of the process
- **Stack:** Temporary data (function parameter, local variables, return addresses etc.)
- Data Section: Global Variables
- **Heap:** dynamically allocated memory during runtime

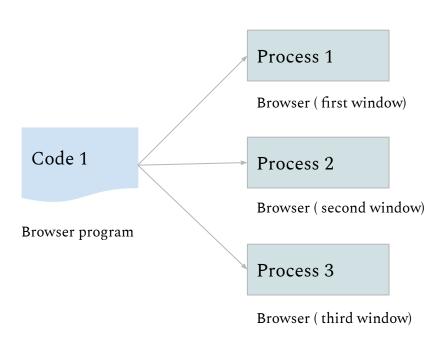


### Program Vs Process

- Program is a collection of instructions that can be executed
- A program is a passive entity.
- A process is an active entity.
- A program becomes a process when it is loaded into memory for execution.



## Same program, Different Process



- Program code is same
- Data, Heap, Stacks contains different information

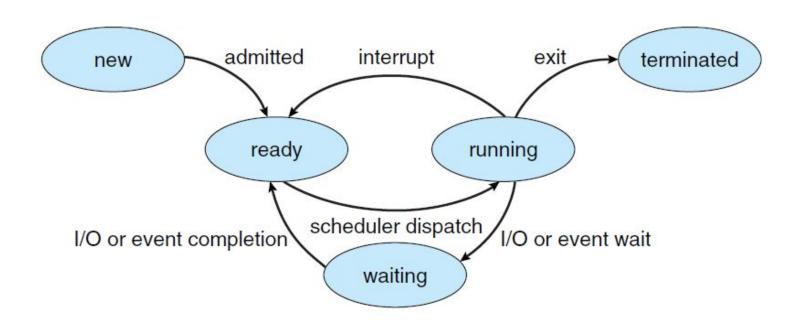
### States of a Process

A process state defines the current activity of that process.

The states a process can be:

- **New**: Process is being created
- **Running:** Instructions are being executed
- Waiting: Process is waiting for some event to occur
- **Ready**: Waiting to be assigned to a processor
- **Terminated**: Process has finished execution

# Process State Diagram



### Representation of Processes in OS

Each process is represented in the operating system by a *Process Control Block* (*PCB*)

PCB is a data structure to store information of Processes such as -

Process state

Program counter

CPU registers

CPU scheduling information

Memory-management information

Accounting information

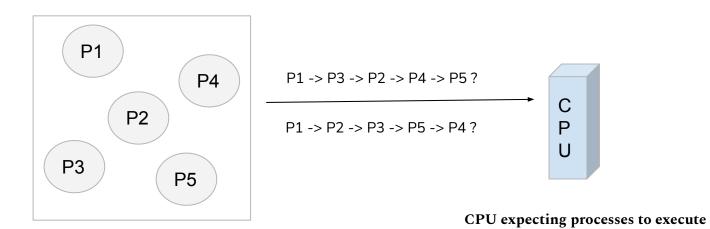
I/O status information

process state process number program counter registers memory limits list of open files

# Operating Systems **Process Scheduling**

# Process Scheduling

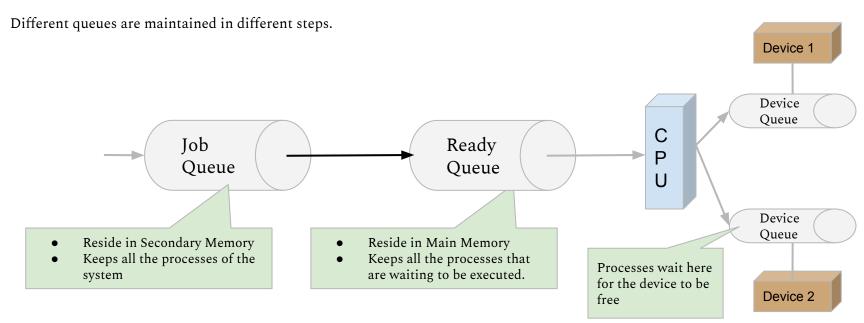
Multiple process is ready to execute. But, which Process should be executed first?



Processes needs to be executed

# Scheduling Queue

Stores the processes in different steps of OS.



# Queueing Diagram

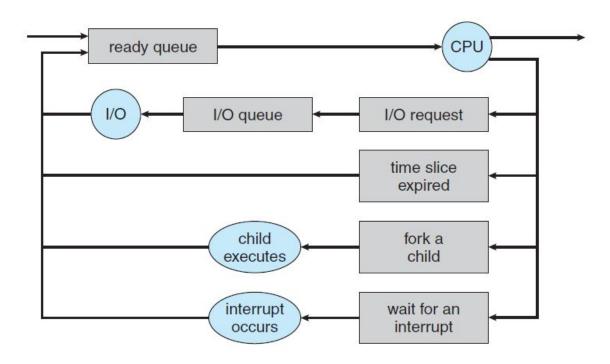
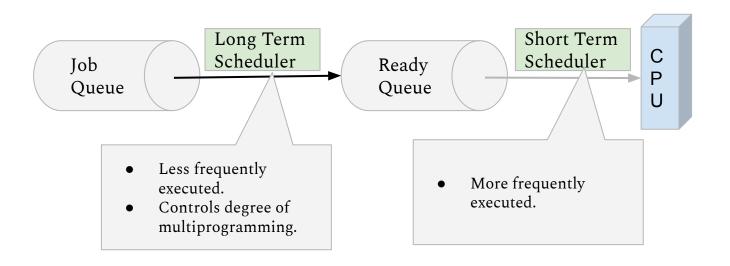


Fig: Representation of Process Scheduling using Queueing-Diagram

#### Schedulers

Schedulers select processes from different queues to be passed to the next phase.

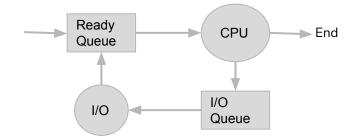


### CPU Bound Vs I/O Bound Process

- CPU bound processes spend more time doing computation using processors than I/O.
- I/O bound processes spend more time in I/O than CPU.

#### Long Term Scheduler must select wisely!

- What will happen if all processes are I/O bound?
- => Empty ready queue
  - What will happen if all processes are CPU bound?



=> Empty waiting queue

#### Medium Term Scheduler

- Time-sharing system may use this scheduler.
- Swapping reduce the degree of multiprogramming.

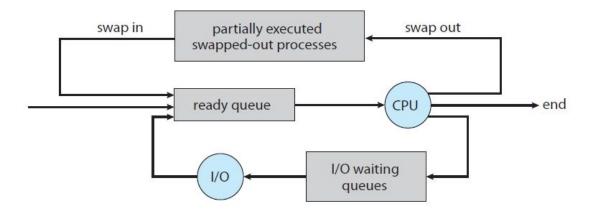


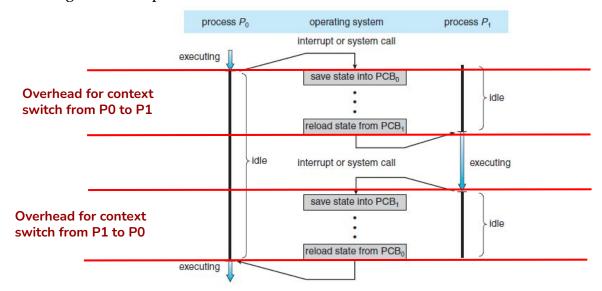
Fig: Addition of swapping in Queueing-Diagram

#### Context Switch

When an interrupt occurs, the system needs to save the current **context** (state) of the process running on the CPU.

Context Switch: 1. Storing currently executed process context

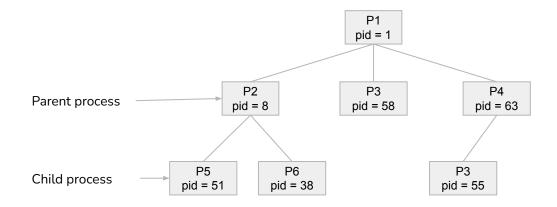
2. Restoring the next process context to execute



# Operating Systems Operations on Process

#### **Process Creation**

- A process is identified by a unique PID (Process Identifier) in the OS.
- A process may create new processes.



- Child process obtain resources from OS or are restricted to Parent's resources
- Parent process may pass initializing data to child process

#### **Process Creation**

• When a process creates new process -

The parent continues to execute concurrently with its children Or,

The parent waits until some or all of its children have terminated

Two address-space possibilities for the new process -

The child process is a duplicate of the parent process
Or
The child process has a new program loaded into it.

#### Process creation in UNIX

System Call: offers the services of the operating system to the user programs.

fork(): create a new process, which becomes the child process of the caller

exec(): runs an executable file, replacing the previous executable

wait(): suspends execution of the current process until one of its children terminates.

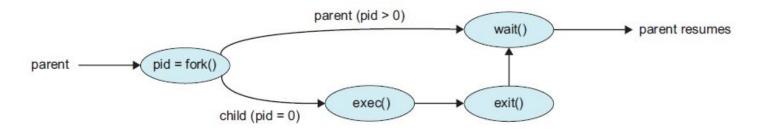


Fig: Process creation using fork() system call

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
pid_t pid;
   /* fork a child process */
   pid = fork();
   if (pid < 0) { /* error occurred */
     fprintf(stderr, "Fork Failed");
     return 1;
   else if (pid == 0) { /* child process */
     execlp("/bin/ls","ls",NULL);
   else { /* parent process */
     /* parent will wait for the child to complete */
      wait(NULL);
     printf("Child Complete");
   return 0;
```

Figure 3.9 Creating a separate process using the UNIX fork() system call.

```
int main(){
   fork();
   fork();
   printf("A");
}
```

```
int main(){
   fork();
   fork();
   fork();
   printf("A");
}
```

```
int main(){
    a = fork();
    if(a==0) fork();
    fork();
    printf("A");
}
```

```
int main(){
    fork();
    a = fork();
    if(a==0) fork();
    printf("A");
}
```

```
int main(){
   int x = 1;
   a = fork();
   if(a==0){
       x = x -1;
       printf("value of x is: %d", x);
   else if (a>0){
       wait(NULL);
       x = x + 1;
       printf("value of x is: %d", x);
```

### Process Termination

A process is terminated when -

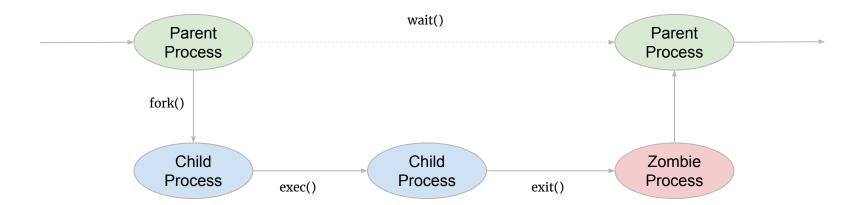
It executes its last statement
Or
Termination cause by another process

When a process is terminated, the resources are deallocated.

A parent may terminate its child if -

- 1. Child has exceeded the usage of resources
- 2. Task assigned to child is no longer needed
- 3. Parent is exiting (cascading termination)

## Zombie Process in UNIX



# Operating Systems Interprocess Communication

## Processes in the system

Processes running concurrently may be -

*Independent* (cannot affect or be affected by other process)

Or

**Cooperating** (can affect or be affected by other process)

Process cooperation is needed for -

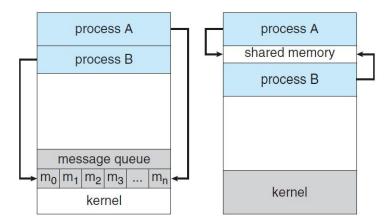
- → Information sharing
- → Computational speedup
- → Modularity
- → Convenience

## Inter Process Communication

IPC is a *mechanism* to exchange data and information among processes.

Two fundamental model of IPC -

- 1. Shared Memory
- 2. Message Passing

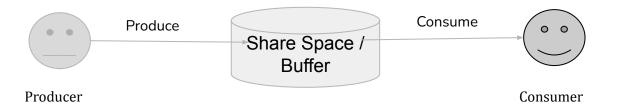


## Shared Memory System

(Producer-Consumer Problem)

Producer: produces products for consumer

Consumer: consumes products provided by producer



#### Producer-Consumer Problem (Producer)

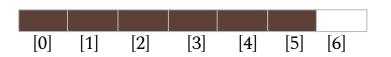
```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

in: next free position in buffer
out: first full position in buffer

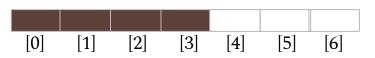
Both initialized with 0.

Here, BUFFER\_SIZE = 7

When buffer is full, in = 6, out = 0



When buffer is not full, In = 4, out = 0



#### Producer-Consumer Problem (Consumer)

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next_consumed */
}
```

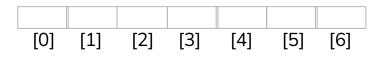
in: next free position in buffer
out: first full position in buffer

Both initialized with 0.

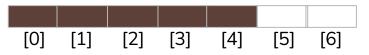
Here, BUFFER SIZE = 7

When buffer is empty,

$$in = 0$$
, out = 0



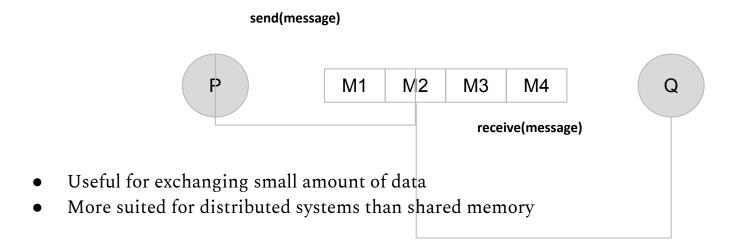
When buffer is not empty, In = 5, out = 0



## Message Passing System

If processes P and Q want to communicate, they must *send* messages to and *receive* messages from each other.

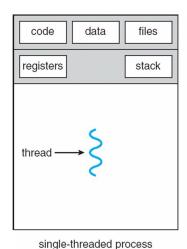
A communication link must exist between P and Q.

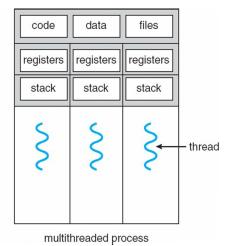


# OPERATING SYSTEMS Threads

### <u>Thread</u>

A thread is a path of execution within a process.





- A thread contains -
  - Thread ID
  - Program Counter
  - Register Set
  - Stack
- Shares with other threads belonging to the same process -
  - Code Section
  - Data Section
  - OS resources

> A traditional process has a single thread of control (Single Threaded Process)

> Process with multiple threads of control, can perform more than one task at a time (Multi Threaded Process)

## **Benefits**

There are four major categories of benefits to multi-threading:

- 1. Responsiveness One thread may provide rapid response while other threads are blocked or slowed down doing intensive calculations.
- 1. Resource sharing By default threads share common code, data, and other resources, which allows multiple tasks to be performed simultaneously in a single address space.
- 1. **Economy** Creating and managing threads ( and context switches between them ) is much faster than performing the same tasks for processes.
- 1. Scalability, i.e. Utilization of multiprocessor architectures A single threaded process can only run on one CPU, no matter how many may be available, whereas the execution of a multi-threaded application may be split amongst available processors. (Note that single threaded processes can still benefit from multi-processor architectures when there are multiple processes contending for the CPU, i.e. when the load average is above some certain threshold.)

## <u>Multicore Programming</u>

**Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:

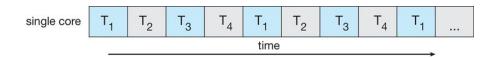
- → Dividing activities
- → Balance
- → Data splitting
- → Data dependency
- → Testing and debugging

Parallelism implies a system can perform more than one task simultaneously

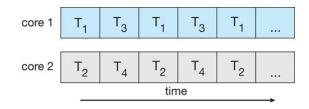
**Concurrency** supports more than one task making progress

 Single processor / core, scheduler providing concurrency

#### Concurrent execution on single-core system:



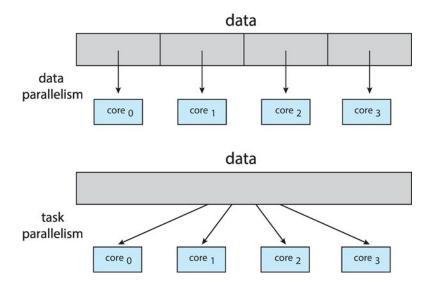
#### Parallelism on a multi-core system:



## Multicore Programming

Data parallelism - distributes subsets of the same data across multiple cores, same operation on each

Task parallelism - distributes threads across cores, each thread performing unique operation



## Amdahl's Law

Identifies performance gains from adding additional cores to an application that has both serial and parallel components

- → S is serial portion
- → N is number of processing cores

$$speedup \le \frac{1}{S + \frac{(1-S)}{N}}$$

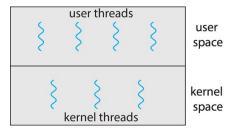
- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches 1 / S

Serial portion of an application has disproportionate effect on performance gained by adding additional cores

# OPERATING SYSTEMS Multithreading Models

### <u>User and Kernel Threads</u>

- There are two types of threads to be managed in a modern system: User threads and kernel threads.
- In a specific implementation, the user threads must be mapped to kernel threads.



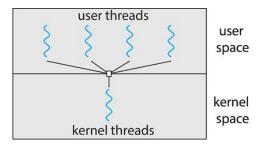
 User threads are supported above the kernel, without kernel support. These are the threads that application programmers would put into their programs.

• Kernel threads are supported within the kernel of the OS itself. All modern OS support kernel level threads, allowing the kernel to perform multiple simultaneous tasks and/or to service multiple kernel system calls simultaneously.

## Multithreading Models

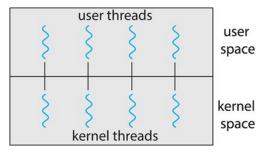
#### Many-to-One:

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Examples:
  - Solaris Green Threads
  - GNU Portable Threads



#### One to One:

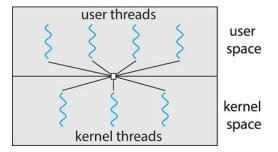
- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows
  - Linux



## Multithreading Models

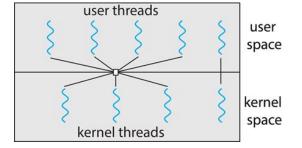
#### Many-to-Many:

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the ThreadFiber package
- Otherwise not very common



#### Two-level Model:

 Similar to M:M, except that it allows a user thread to be bound to kernel thread



# OPERATING SYSTEMS Thread Libraries

## Thread library

- Thread libraries provide programmers with an API for creating and managing threads.
- Thread libraries may be implemented either in user space or in kernel space. The former
  involves API functions implemented solely within user space, with no kernel support. The latter
  involves system calls, and requires a kernel with thread library support.
- There are three main thread libraries in use today:
  - POSIX Pthreads may be provided as either a user or kernel library, as an extension to the POSIX standard.
  - $\circ$  Win32 threads provided as a kernel-level library on Windows systems.
  - Java threads Since Java generally runs on a Java Virtual Machine, the implementation of threads is based upon whatever OS and hardware the JVM is running on, i.e. either Pthreads or Win32 threads depending on the system.

### **Pthreads**

- The POSIX standard (IEEE 1003.1c) defines the specification for pThreads, not the implementation.
- pThreads are available on Solaris, Linux, Mac OSX, Tru64, and via public domain shareware for Windows.
- Global variables are shared amongst all threads.
- One thread can wait for the others to rejoin before continuing.
- pThreads begin execution in a specified function.
- Common in UNIX operating systems (Linux & Mac OS X)

## <u> Iava Threads</u>

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by:
  - Extending Thread class
  - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

→ Standard practice is to implement Runnable interface

# OPERATING SYSTEMS Threading Issues

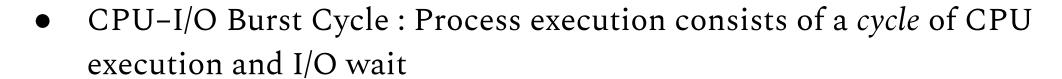
## Threading Issues

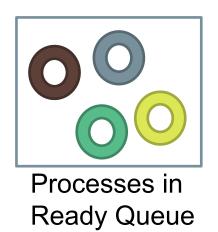
- fork() and exec() System Calls: Duplicate all the threads or not?
- Thread cancellation: Thread cancellation is the task of terminating a thread before it has completed.
- Signal Handling: Where should a signal be delivered?
- Thread Pool: Create a number of threads at the process start-up.
- Thread Specific data: Each thread might need it's own copy of certain data.

# OPERATING SYSTEMS CPU Scheduling

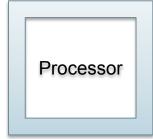
## Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- Continuous Cycle:
  - one process has to wait (I/O)
  - Operating system takes the CPU away
  - Give CPU to another process
  - This pattern continues









## CPU Scheduler

- Selects from among the processes in ready queue, and allocates the CPU to one of them
  - FIFO queue
  - Priority queue
  - Tree
  - Unordered linked-list
- CPU scheduling decisions may take place when a process:
  - 1. Switches from running to waiting state (I/O request)
  - 2. Switches from running to ready state (e.g. when interrupt occurs)
  - 3. Switches from waiting to ready (e.g. at completion of I/O)
  - 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is *preemptive* 
  - Consider access to shared data
  - Consider preemption while in kernel mode
  - Consider interrupts occurring during crucial OS activities

## Scheduling Criteria

- **CPU utilization** keep the CPU as busy as possible
- Throughput # of processes that complete their execution per time unit

#### • Turnaround time

- amount of time to execute a particular process
- -- the interval from the time of submission of a process to the time of the completion.
- -- sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, doing I/O
- Waiting time amount of time a process has been waiting in the ready queue
- **Response time** amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

## Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

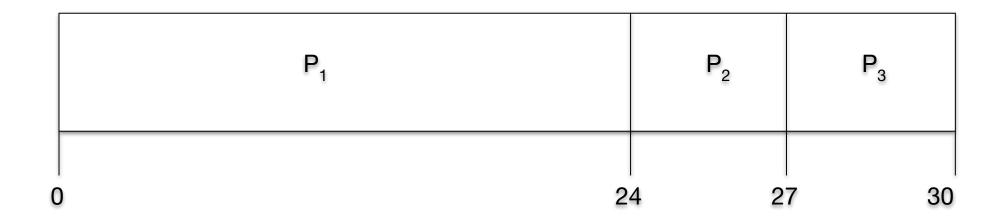
### OPERATING SYSTEMS

# CPU Scheduling Algorithms - First Come First Serve (FCFS)

## First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
$P_{1}$	24
$P_{2}$	3
$P_{3}$	3

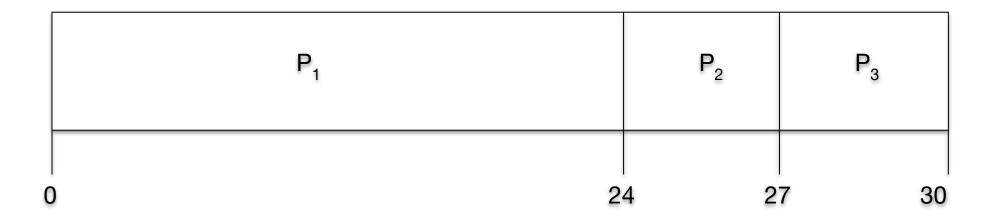
• Suppose that the processes arrive in the order:  $P_1$ ,  $P_2$ ,  $P_3$  The Gantt Chart for the schedule is:



## First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
$P_{1}$	24
$P_{2}$	3
$P_{3}$	3

• Suppose that the processes arrive in the order:  $P_1$ ,  $P_2$ ,  $P_3$  The Gantt Chart for the schedule is:



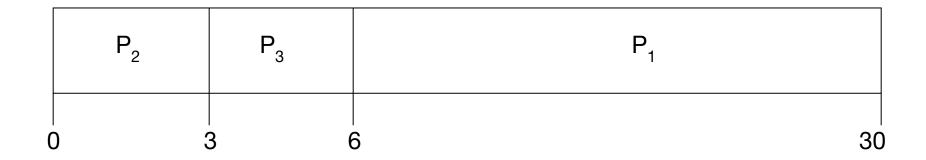
- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time: (0 + 24 + 27)/3 = 17
- Turnaround time  $P_1 = 24$ ;  $P_2 = 27$ ;  $P_3 = 30$

## FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

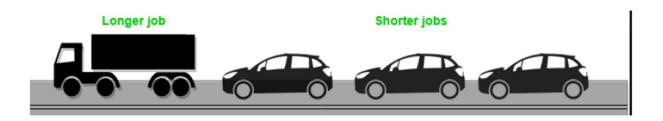
• The Gantt chart for the schedule is:



Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ,  $P_3 = 3$ Average waiting time: (6 + 0 + 3)/3 = 3Much better than previous case

**Convoy effect** - short process behind long process.

Consider one CPU-bound and many I/O-bound processes



# CPU Scheduling Algorithms - Shortest Job First (SJF)

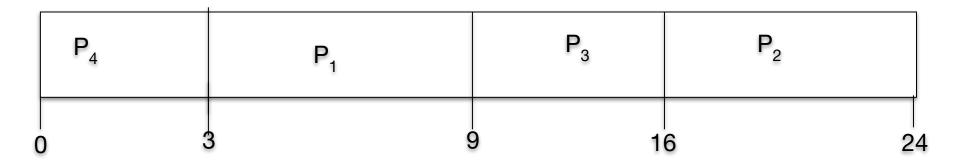
### Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - Non-preemptive once CPU given to the process it cannot be preempted until completes its CPU burst
  - preemptive if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal gives minimum average waiting time for a given set of processes

### Example of SJF

# ProcessBurst Time $P_1$ 6 $P_2$ 8 $P_3$ 7 $P_4$ 3

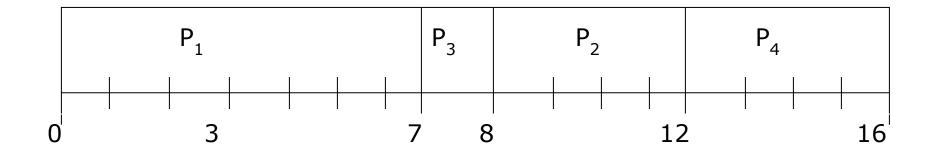
• SJF scheduling chart



### Example of Non-Preemptive SJF

<u>Process</u>	<u> Arrival Time</u>		<u>Burst Time</u>
$P_{1}$	0.0	7	
$P_{2}$	2.0	4	
$P_3$	4.0	1	
$P_{_{4}}$	5.0	4	

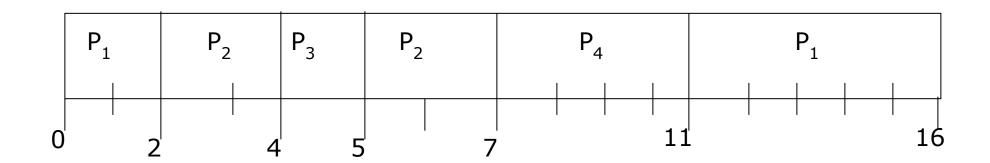
• SJF (non-preemptive)



### Example of Preemptive SJF

<u>Process</u>	<u> Arrival Time</u>		Burst Time	
$P_{1}$	0.0	7		
$P_{2}^{}$	2.0	4		
$P_{3}$	4.0	1		
$P_{_{4}}$	5.0	4		

• SJF (preemptive)

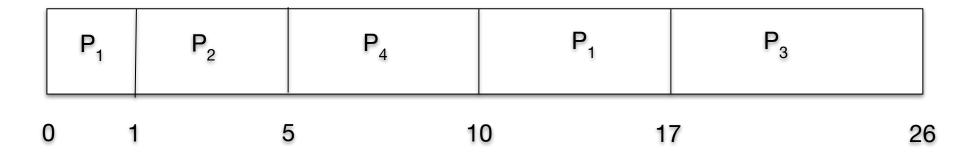


### Example of Shortest-remaining-time-first

• Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	8
P2	1	4
P3	2	9
P4	3	5

• Preemptive SJF Gantt Chart



## CPU Scheduling Algorithms - Priority Scheduling

### Priority Scheduling

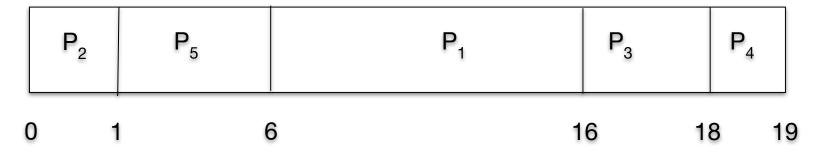
- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer ≡ highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Priority can be defined either internally or externally.
  - Factors for internal priority assignment:
    - Time limit, memory requirements, the number or open files etc.
  - Factors for external priority assignment:
    - Importance of the process, the type and amount of funds being paid for computer use, department sponsoring works etc.

### Example of Priority Scheduling

#### Non Preemptive:

<u>Process</u>	<u>Bu</u>	<u>rst Time</u>	<u>Priority</u>
$P_{_{1}}$	10	3	
$P_{2}$	1	1	
$P_{3}$	2	4	
$P_{\underline{4}}$	1	5	
$P_{\overline{5}}$	5	2	

Priority scheduling Gantt Chart



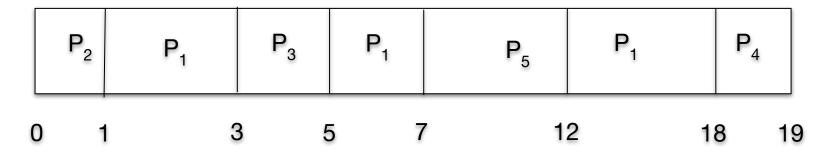
- Problem ≡ **Starvation** low priority processes may never execute
- Solution  $\equiv$  **Aging** as time progresses increase the priority of the process

### Example of Priority Scheduling

#### **Preemptive**:

<u>Process</u>	Arrival '	<u> </u>	<u>Burst Time</u>	<b>Priority</b>
$P_{_{1}}$	0	10	4	
$P_{2}$	0	1	1	
$P_{3}$	3	2	3	
$P_{_{4}}$	5	1	5	
$P_{5}$	7	5	2	

Priority scheduling Gantt Chart



- Problem ≡ **Starvation** low priority processes may never execute
- Solution  $\equiv$  **Aging** as time progresses increase the priority of the process

## CPU Scheduling Algorithms - Round Robin (RR)

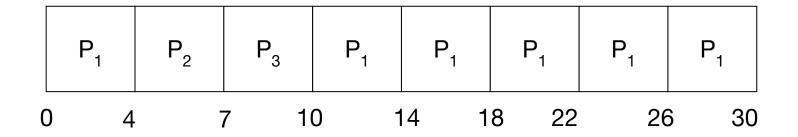
### Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q, then each process gets 1/n of the CPU time in chunks of at most q time units at once. No process waits more than (n-1)q time units.
- Timer interrupts every quantum to schedule next process
- Performance
  - $\circ$  q large  $\Rightarrow$  FIFO
  - o  $q \text{ small} \Rightarrow q \text{ must be large with respect to context switch, otherwise overhead is too high}$

### Example of RR with Time Quantum = 4

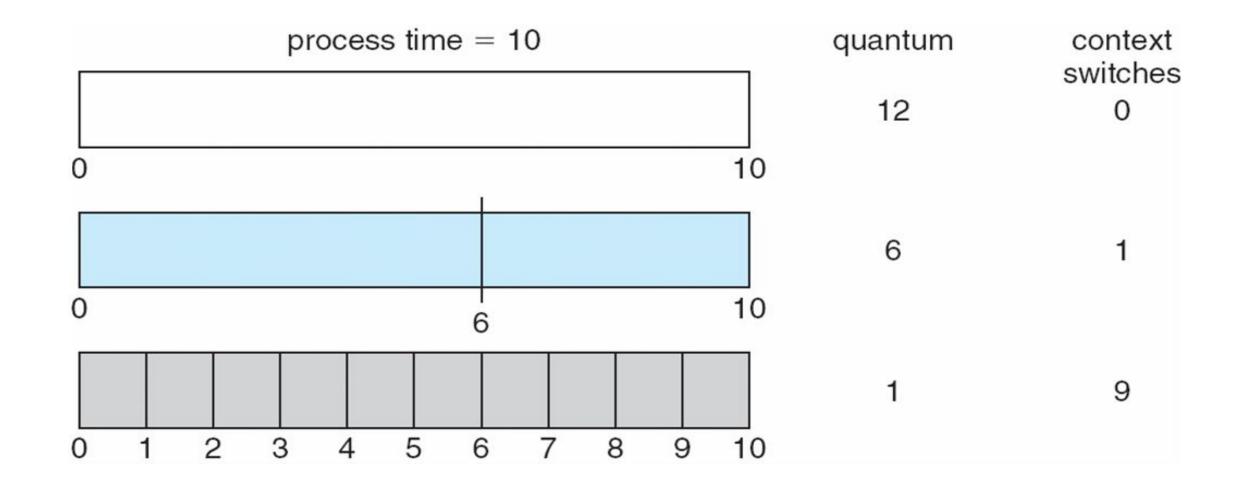
<u>Process</u>	<b>Burst Time</b>
$P_{_1}$	24
$\vec{P}_2$	3
$P_{3}^{2}$	3
3	

The Gantt chart is:

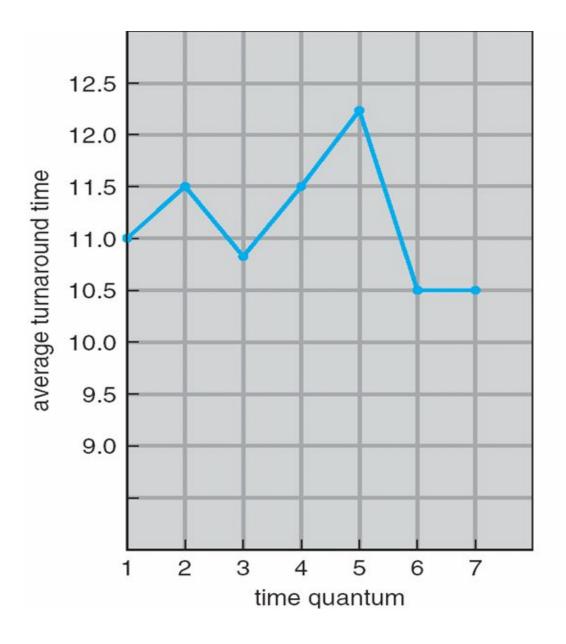


- Average waiting time is 17/3 = 5.66 milisecond
- Typically, higher average turnaround than SJF, but better response
- quantum should be large compared to context switch time
- Quantum usually 10ms to 100ms, context switch < 10 usec

### Time Quantum and Context Switch Time



### Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

80% of CPU bursts should be shorter than quantum

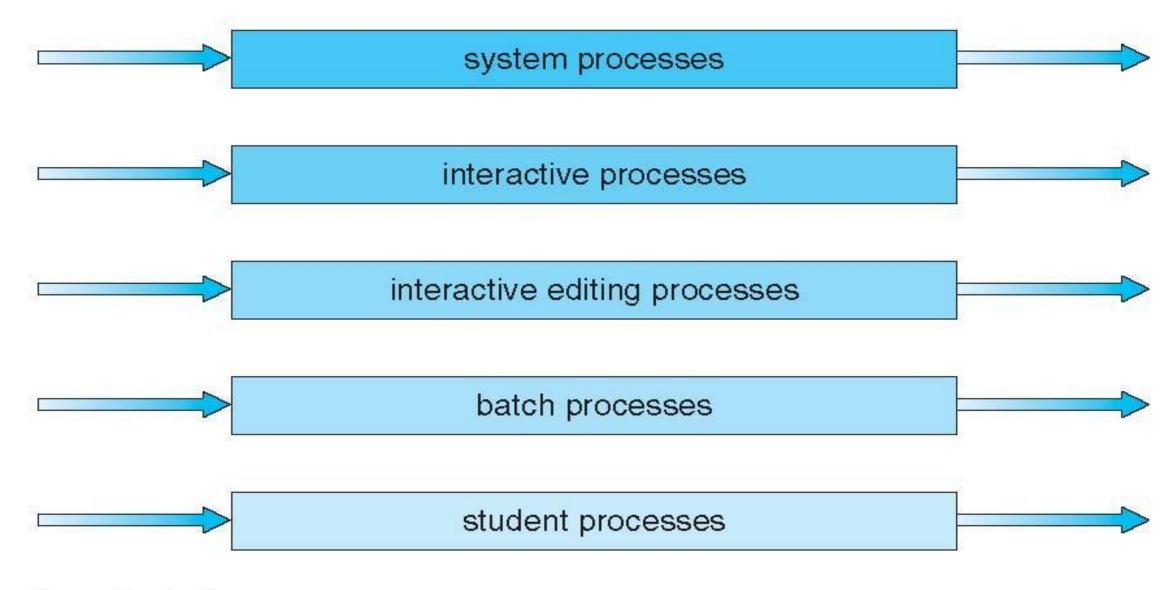
# CPU Scheduling Multilevel Queue, Multilevel Feedback Queue

### Multilevel Queue

- Another class of scheduling algorithm needs- in which processes are classified into different groups, e.g.:
  - o foreground (interactive) processes
  - background (batch) processes
- They have different response time requirements-so different scheduling needs.
- Foreground processes may have priority over background processes.
- A multilevel queue-scheduling algorithm partitions the ready queue into several separate queues-we
  can see it in the figure of next slide:-
- Each queue has its own scheduling algorithm:
  - Foreground queue scheduled by RR algorithm
  - Background queue scheduled by FCFS algorithm
- Scheduling must be done between the queues:
  - Fixed priority preemptive scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - Time slice each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., foreground queue can be given 80% of the CPU time for RR-scheduling among its processes, while 20% to background in FCFS manner.

### Multilevel Queue Scheduling

highest priority



lowest priority

### Multilevel Feedback Queue scheduling

- Multilevel Feedback Queue scheduling, allows a process to move between queues.
- If a process uses too much CPU time, it will be moved to a lower priority queue.
- Similarly, a process that waits too long in a lower-priority queue may me moved to a higher-priority queue.
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - o scheduling algorithms for each queue
  - o method used to determine when to upgrade a process
  - o method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

### Example of Multilevel Feedback Queue

- Three queues: (can see the figure in next slide)
  - $\circ$   $Q_0$  RR with time quantum 8 milliseconds
  - $\circ$   $Q_1$  RR time quantum 16 milliseconds
  - $\circ$   $Q_2 FCFS$
- Scheduling
  - A new job enters queue  $Q_0$  which is served for RR
    - When it gains CPU, job receives 8 milliseconds
    - If it does not finish in 8 milliseconds, job is moved to queue Q<sub>1</sub>
  - At  $Q_1$  job is again served RR and receives 16 additional milliseconds
    - If it still does not complete, it is preempted and moved to queue  $Q_2$

